

# CSCI567 Machine Learning (Spring 2018)

Michael Shindler

Lecture 24: April 18, 2018

# Outline

- 1 Administration
- 2 Recommender Systems: Introduction
- 3 Content-based recommendations
- 4 Collaborative Filtering
- 5 Hybrids

# Outline

- 1 Administration
- 2 Recommender Systems: Introduction
- 3 Content-based recommendations
- 4 Collaborative Filtering
- 5 Hybrids

# Administrative

- Quiz 3 coming up
- Remember your ID, pencil(s)

# Outline

1 Administration

2 Recommender Systems: Introduction  
• History

3 Content-based recommendations

4 Collaborative Filtering

5 Hybrids

# Example: Amazon.com

Deals recommended for you [See all deals](#)



\$105.99 - \$1,289.99  
Ends in 03:52:40



\$48.95  
\$69.99  
Ends in 03:42:41



\$84.99 - \$104.99  
\$20.35  
Ends in 52:41



\$12.7C  
Ends in C

# Example: The Facebook

Suggested Pages [See All](#)

351 people like this.

 **XFL Access**  
[Website](#)



 [Like Page](#)

# Example: The Facebook

 **Star Trek: Discovery**  
Sponsored • 

Your journey begins NOW! Stream season one of Star Trek: Discovery across devices.



Try 1 Week FREE  
Exclusively on CBS All Access  
[CBS.COM/STAR-TREK-DISCOVERY/](http://CBS.COM/STAR-TREK-DISCOVERY/) [Sign Up](#)

 680      84 Comments 29 Shares 13K Views

 Like       Comment       Share

# Example: Netflix

Because you added Heat to your list



# Example: Netflix

Because you watched Star Trek: Deep Space Nine



# Example: Movie recommendation

Training data

user	movie	score
1	21	1
1	213	5
2	345	4
2	123	4
2	768	3
3	76	5
4	45	4
5	568	1
5	342	2
5	234	2
6	76	5
6	56	4

Test data

user	movie	score
1	62	?
1	96	?
2	7	?
2	3	?
3	47	?
3	15	?
4	41	?
4	28	?
5	93	?
5	74	?
6	69	?
6	83	?

# Recommender Systems: Other Examples

- Recommendation of groups, jobs or people on LinkedIn
- Friend recommendation and ad personalization on Facebook
- News recommendation at Forbes.com
- etc.

**LinkedIn Jobs**

Parker,

Check out these jobs that may interest you:

	<b>Director, Product Management</b> Move, Inc. - San Francisco Bay Area	<a href="#">View Job &gt;</a>
	<b>Director - Product Management</b> Symantec - San Francisco Bay Area	<a href="#">View Job &gt;</a>
	<b>Director of Product Management, Mobile Games</b> iWin, Inc. - San Francisco Bay Area	<a href="#">View Job &gt;</a>
	<b>Director Product Management, Mobile Carriers</b> Ruckus Wireless - San Francisco Bay Area	<a href="#">View Job &gt;</a>
	<b>Director of Product Management - eCommerce</b> youSendIt - San Francisco Bay Area	<a href="#">View Job &gt;</a>

[See more jobs you may be interested in >](#)

# Why using Recommender Systems?

- Value for the customer
  - Find things that are interesting
  - Narrow down the set of choices
  - Help me explore the space of options
  - Discover new things
  - Entertainment, ...
- Value for the provider
  - Additional and probably unique personalized service for the customer
  - Increase sales, click trough rates, conversion etc.
  - Opportunities for promotion, persuasion
  - Obtain more knowledge about customers

# Value of Recommender Systems

- Netflix: 2/3 of the movies watched are recommended
- Google News: recommendations generate 38% more clickthrough
- Amazon: 35% sales from recommendations
- Choicestream: 28% of the people would buy more music if they found what they liked.

# Recommender Systems: Problem

Estimate a utility function that automatically predicts how a user will like an item.

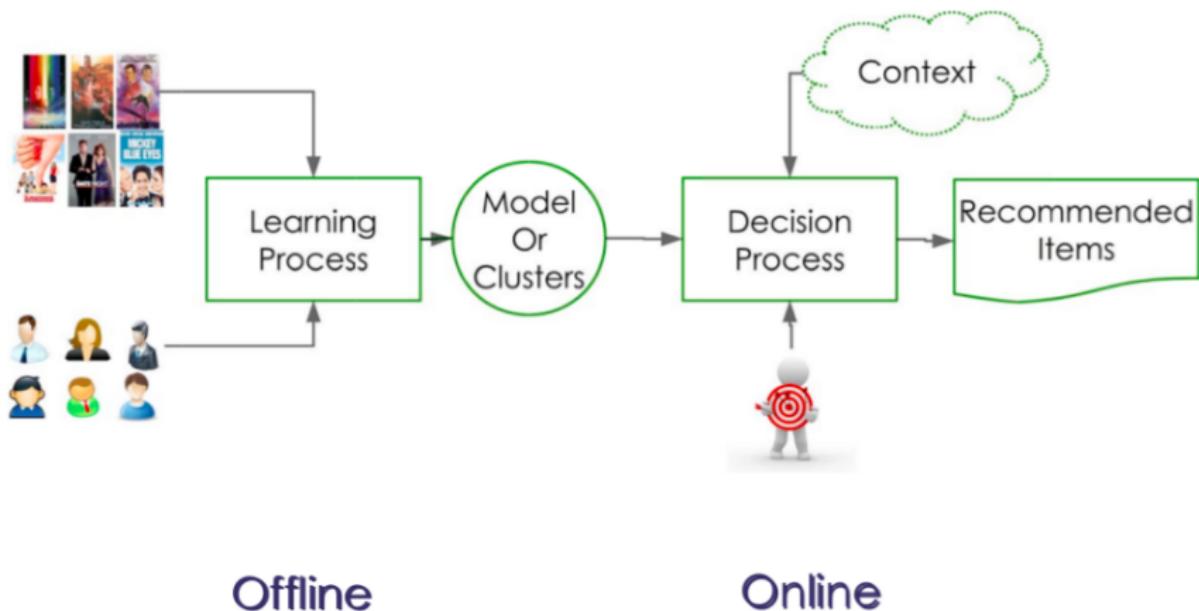
Based on:

- Past behavior
- Relations to other users
- Item similarity
- Context
- ...

# Recommender Systems: Problem

- Let  $U$  be set of all users and let  $I$  be set of all possible recommendable items
- Let  $F$  be a utility function measuring the usefulness of item  $i$  to user  $u$ , i.e.,  $F : I \times U \rightarrow R$ , where  $R$  is a totally ordered set.
- For each user  $u \in U$ , we want to choose items  $i \in I$  that maximize  $F$ .

# Recommender Systems: Two-step Process

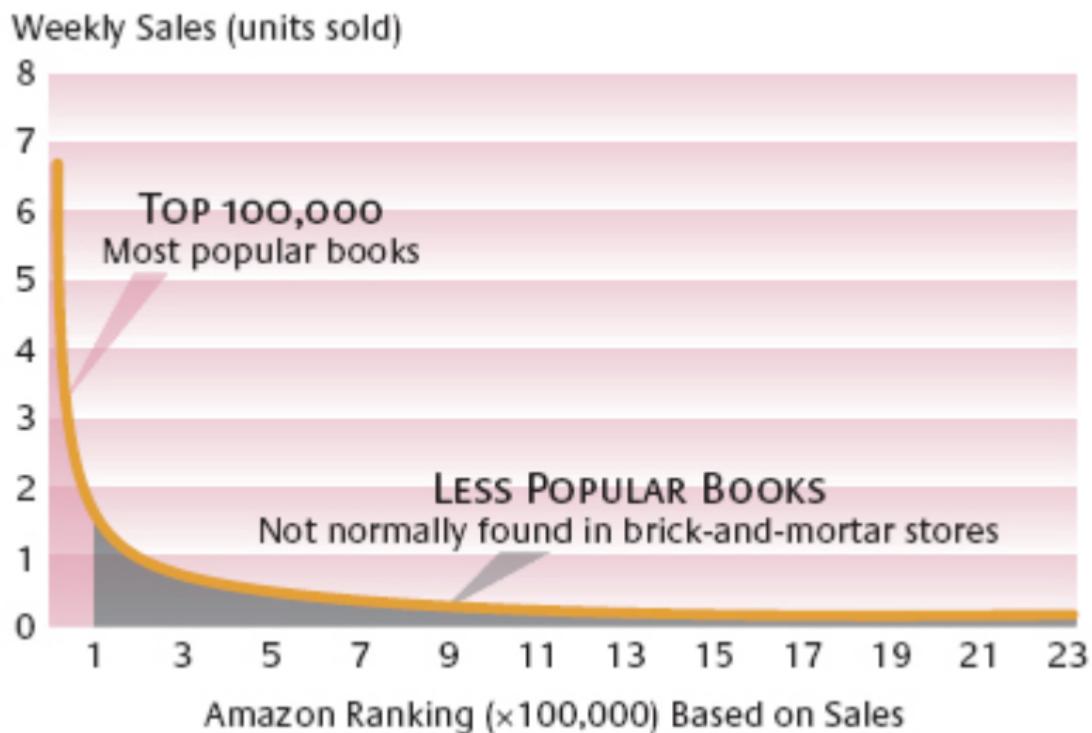


# Recommender Systems: Approaches

Popular approaches:

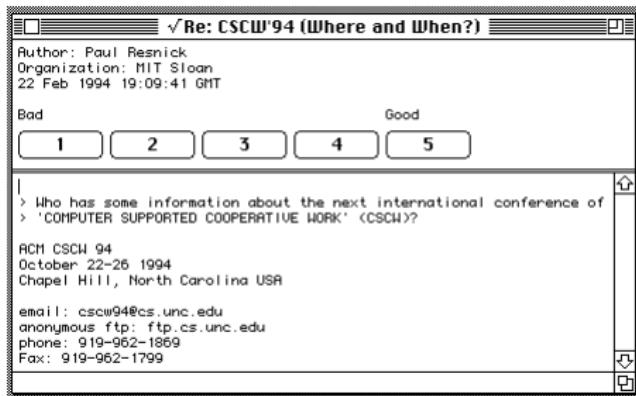
- Content-based recommendations:
  - The user will be recommended items based on profile information or similar to the ones the user preferred in the past
- Collaborative filtering (or collaborative recommendations):
  - The user will be recommended items that people with similar tastes and preferences liked in the past
- Hybrids:
  - Combine collaborative and content-based methods

# Why so many online sources?



# The first collaborative filtering

- Tapestry, 1994:  
<http://ccs.mit.edu/papers/CCSWP165.html>



# The Netflix Challenge

- <https://www.netflixprize.com>
- October 2006: Netflix Prize for improving recommendation algorithm.
- \$1 Million to whoever gets 10% improvement.
- Netflix published a dataset:
  - $\approx 17,000$  movies
  - 500,000 users.
  - 100 million ratings (training set was 99 million)
- Judged on root mean-square error

# The Netflix Challenge

- Lasted about three years
- Eventually won: “Bellkor’s Pragmatic Chaos” at 10.6%
- 40,000 teams entered, 186 countries

# Outline

- 1 Administration
- 2 Recommender Systems: Introduction
- 3 Content-based recommendations
- 4 Collaborative Filtering
- 5 Hybrids

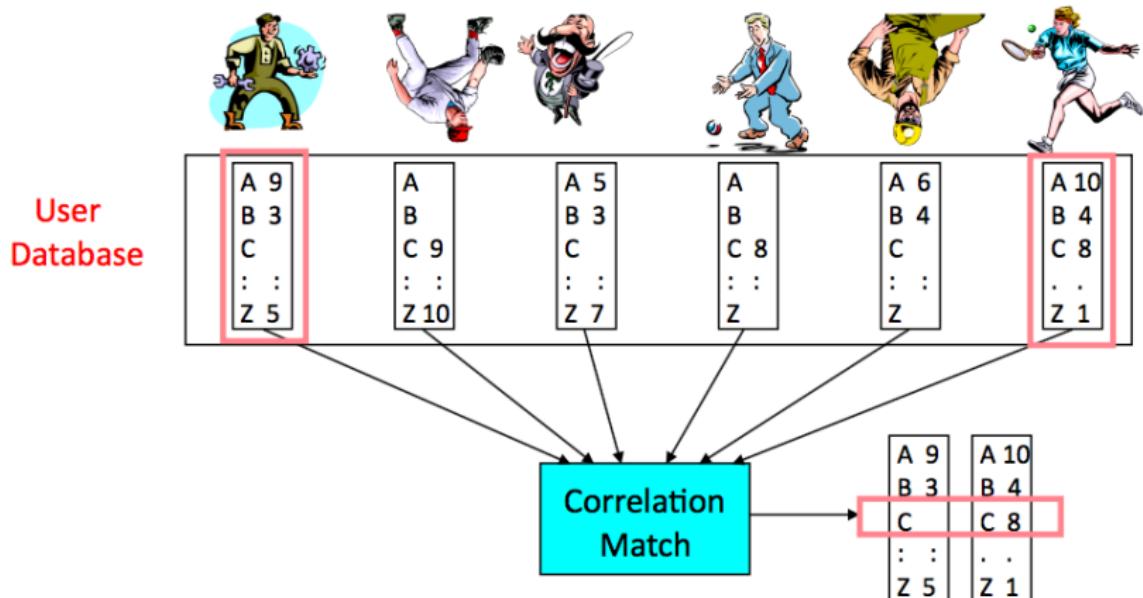
# Content-based recommendations

- Recommend items that matches the User Profile.
- The Profile is based on items user has liked in the past or explicit interests that he defines.
- A content-based recommender system matches the profile of the item to the user profile to decide on its relevancy to the user.

# Outline

- 1 Administration
- 2 Recommender Systems: Introduction
- 3 Content-based recommendations
- 4 Collaborative Filtering
- 5 Hybrids

# Collaborative Filtering



# Collaborative Filtering

- Collaborative filtering (CF):
  - k-nearest neighbor (kNN)
  - matrix factorization
- The most prominent approach to generate recommendations: Used by large, commercial e-commerce sites
- Well-understood, various algorithms and variations exist
- Applicable in many domains (book, movies, DVDs, ..)
- Approach: use the wisdom of the crowd to recommend items
- Basic assumption and idea
  - Users give ratings to catalog items (implicitly or explicitly)
  - Customers who had similar tastes in the past, will have similar tastes in the future

# Collaborative Filtering: User-based Model

GroupLens: an open architecture for collaborative filtering of netnews, P. Resnick et al., ACM CSCW, 1994.

- Goal: given an active user (Alice) and an item  $i$  not yet seen by Alice, the goal is to estimate Alices rating for this item.
- Basic idea: People who agreed in their subjective evaluations in the past are likely to agree again in the future
- Algorithm: find a set of users (peers) who liked the same items as Alice in the past and who have rated item  $p$ , e.g. the average of their ratings to predict, if Alice will like item  $p$ ; do this for all items (Alice has not seen and recommend the best-rated)

	Item1	Item2	Item3	Item4	Item5
Alice	5	3	4	4	?
User1	3	1	2	3	3
User2	4	3	4	3	5
User3	3	3	1	5	4
User4	1	5	5	2	1

# Collaborative Filtering: Important Questions

- How do we measure similarity?
- How many neighbors should we consider?
- How do we generate a prediction from the neighbors ratings?

# Collaborative Filtering: Important Questions

- How do we measure similarity?

Answer: Pearson correlation

$$\text{sim}(a, b) = \frac{\sum_p (r_{a,p} - \bar{r}_a)(r_{b,p} - \bar{r}_b)}{\sqrt{\sum_p (r_{a,p} - \bar{r}_a)^2} \sqrt{\sum_p (r_{b,p} - \bar{r}_b)^2}}$$

- How many neighbors should we consider?

Answer: Use similarity threshold or fixed number of neighbors.

# Collaborative Filtering: Important Questions

- How do we generate a prediction from the neighbors ratings?

Answer:

$$\text{pred}(a, p) = \bar{r}_a + \frac{\sum_{b \in N(a)} \text{sim}(a, b) * (r_{b,p} - \bar{r}_b)}{\sum_{b \in N(a)} \text{sim}(a, b)}$$

- Calculate whether the neighbors ratings for the unseen item  $i$  are higher or lower than their average
- Combine the rating differences using the similarity as a weight
- Add/subtract the neighbors bias from the active users average and use this as a prediction

# User-based CF: Example



	SHERLOCK	HOUSE OF CARDS	THE AVENGERS	COMMUNITY	BREAKING BAD	THE WALKING DEAD	sim(u,v)
1	2			4	5		NA
2			4			1	NA
3			5		2		
4		1		5		4	
5			4			2	
6	4	5		1			

# User-based CF: Example

sim(u,v)

	SHERLOCK	HOUSE OF CARDS	THE AVENGERS	A FRESH PRINCE OF BEL-AIR	BREAKING BAD	THE WALKING DEAD	
2	2			4	5		NA
5	5		4			1	0.87
			5		2		
		1		5		4	
			4			2	
4	4	5		1			NA



# User-based CF: Example



# User-based CF: Example

sim(u,v)

							sim(u,v)
	2			4	5		NA
	5		4			1	0.87
			5		2		1
		1		5		4	-1
			4			2	NA
	4	5		1			

# User-based CF: Example



A user-based collaborative filtering matrix showing ratings for various TV shows across different users. The matrix is 6 rows by 6 columns. Rows represent users and columns represent TV shows. The matrix is partially filled with numerical ratings (2, 4, 5) and some entries are marked with an asterisk (\*). The last row contains all NA values.

	SHERLOCK	HOUSE OF CARDS	THE AVENGERS	A MIDDLE-AGED ADULT DEVELOPMENT	BREAKING BAD	THE WALKING DEAD	sim(u,v)
User 1	2			4	5		NA
User 2	5		4			1	0.87
User 3			5		2		1
User 4			1		5		-1
User 5	3.51*	3.81*	4	2.42*	2.48*	2	
User 6	4	5		1			NA

## User-based CF: Challenges

- Sparsity evaluation of large item sets, users purchases are under 1%.
- Difficult to make predictions based on nearest neighbor algorithms → Accuracy of recommendation may be poor.
- Scalability - Nearest neighbor require computation that grows with both the number of users and the number of items.
- Poor relationship among like minded but sparse-rating users.
- Solution : usage of latent models to capture similarity between users and items in a reduced dimensional space

# Collaborative Filtering: Item-based Model

Item-based collaborative filtering recommendation algorithms, B. Sarwar et al., WWW 2001.

- Scalability issues arise with user-based model if there are many more users than items ( $m \gg n$ ,  $m = \#$  of users,  $n = \#$  of items). e.g. amazon.com.
  - Space complexity is  $O(m^2)$  when pre-computed;
  - Time complexity for computing Pearson is  $O(m^2n)$ .
- High sparsity leads to few common ratings between two users
- Basic idea: Item-based CF exploits relationships between items first, instead of relationships between users.

	Item1	Item2	Item3	Item4	Item5
Alice	5	3	4	4	?
User1	3	1	2	3	3
User2	4	3	4	3	5
User3	3	3	1	5	4
User4	1	5	5	2	1

# Collaborative Filtering: Item-based Model

- Item-based filtering does not solve the scalability problem itself
- Pre-processing approach by Amazon.com (in 2003)
  - (1) Calculate all pair-wise item similarities in advance
  - (2) The neighborhood to be used at run-time is typically rather small, because only items are taken into account which the user has rated
  - (3) Item similarities are supposed to be more stable than user similarities
- Memory requirements: Up to  $N^2$  pair-wise similarities to be memorized ( $N$  = number of items) in theory; In practice, this is significantly lower (items with no co-ratings)
- Minimum threshold for co-ratings (items, which are rated at least by  $n$  users)
- Limit the size of the neighborhood (might affect recommendation accuracy)

# Item-based CF: Example



# Item-based CF: Example



# Item-based CF: Example



# Item-based CF: Example



# Item-based CF: Example



# Item-based CF: Example



	2			4	5	2.94*	
	5		4			1	
			5		2	2.48*	
		1		5		4	
			4			2	
	4	5		1		1.12*	
	sim(i,j)	-1	-1	0.86	1	NA	

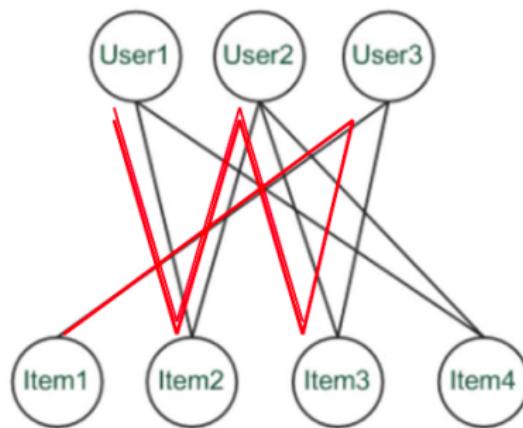
# Performance Implications

- Bottleneck - Similarity computation.
- Time complexity, highly time consuming with millions of users and items in the database.
  - Isolate the neighborhood generation and predication steps.
  - off-line component / model similarity computation, done earlier and stored in memory.
  - on-line component prediction generation process.

# Collaborative Filtering: Graph-based Model

Spreading activation (sketch)

- Use paths of lengths  $\geq 3$  to recommend items
- Length 3: Recommend Item 3 to User 1
- Length 5: Item 1 also recommendable



# Collaborative Filtering: Dimension Reduction

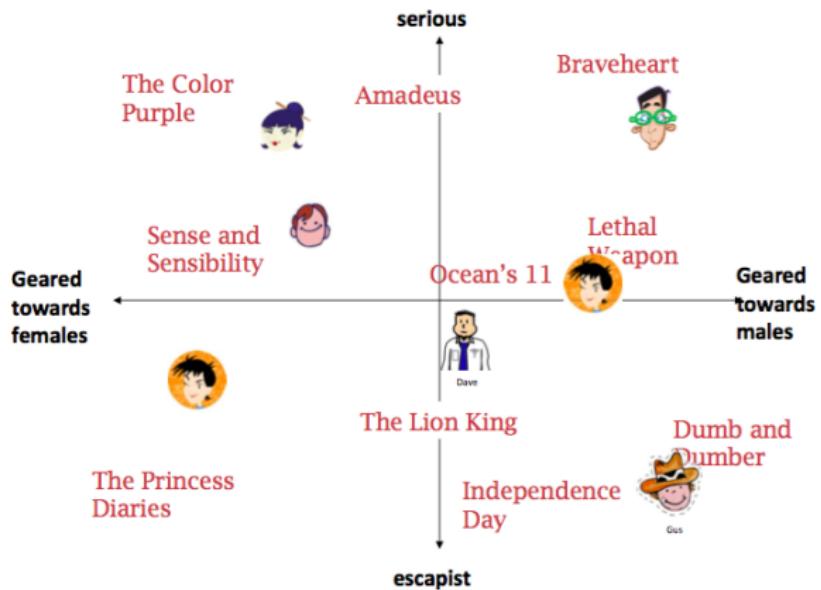
Application of Dimensionality Reduction in Recommender System, B. Sarwar et al., WebKDD Workshop, 2000.

Basic idea: Trade more complex offline model building for faster online prediction generation

- Singular Value Decomposition for dimensionality reduction of rating matrices
- Captures important factors/aspects and their weights in the data factors can be genre, actors, or some semantics
- Assumption that  $k$  dimensions capture the signals and filter out noise

# Collaborative Filtering: Dimension Reduction

Application of Dimensionality Reduction in Recommender System, B. Sarwar et al., WebKDD Workshop, 2000.



# Collaborative Filtering: Dimension Reduction

Application of Dimensionality Reduction in Recommender System, B. Sarwar et al., WebKDD Workshop, 2000.

- SVD:  $M_k = U_k \times \Sigma_k \times V_k^T$

$U_k$	Dim1	Dim2
Alice	0.47	-0.30
Bob	-0.44	0.23
Mary	0.70	-0.06
Sue	0.31	0.93

$V_k^T$							
Dim1	-0.44	-0.57	0.06	0.38	0.57		
Dim2	0.58	-0.66	0.26	0.18	-0.36		

$\Sigma_k$	Dim1	Dim2
Dim1	5.63	0
Dim2	0	3.23

- Prediction:  $\hat{r}_{ui} = \bar{r}_u + U_k(Alice) \times \Sigma_k \times V_k^T(EPL)$   
 $= 3 + 0.84 = 3.84$

# Collaborative Filtering: Latent Class Models

Latent Class Models for Collaborative Filtering, Hofmann and Puzieha, IJCAI 1999.

- Input: user  $u$  and item  $i$  with rating  $v$
- One example of the latent class model: the vote on an item is independent of the users identity given the true value of the latent variable  $z$ , i.e.,

$$P(v|u, i) = \sum_z P(v|u, i, z)P(z|u) = \sum_z P(v|i, z)P(z|u).$$

# Collaborative Filtering: Hybrid Approach

Factorization meets the neighborhood: a multifaceted collaborative filtering model, Y. Koren, ACM SIGKDD, 2008.

Basic idea: Combination of different approaches in one prediction single function

- Baseline estimates:  $v_{u,i} = \mu + b_u + b_i$ . The resulting optimization problem is:

$$\min_{b^*} \sum_{(u,i) \in \mathcal{K}} (v_{u,i} - \mu - b_u - b_i)^2 + \lambda_1 (\sum_{u'} b_{u'}^2 + \sum_{v'} b_{v'}^2)$$

- Baseline + Latent model:  $v_{u,i} = \mu + b_u + b_i + p_u^T q_i$ . The resulting optimization problem is:

$$\begin{aligned} \min_{b^*, p^*, q^*} & \sum_{(u,i) \in \mathcal{K}} (v_{u,i} - \mu - b_u - b_i - p_u^T q_i)^2 + \lambda_1 (\sum_{u'} b_{u'}^2 + \sum_{v'} b_{v'}^2) \\ & + \lambda_2 (||p||_2 + ||q||_2) \end{aligned}$$

# Evaluation Metrics

- Novelty: The ability of a recommender system to recommend items that the user was not already aware of.
- Serendipity: Users are given recommendations for items that they would not have seen given their existing channels of discovery. A measure of how surprising the recommendations are
- Coverage: The percentage of the items known to the recommender system for which it can generate predictions.
- Learning Rate: How quickly the CF system becomes an effective predictor of taste as data begins to arrive
- Confidence: Ability to evaluate the likely quality of its predictions
- User Satisfaction: By surveying the users or measuring retention and use statistics

# Evaluation Metrics

- Accuracy
  - Predict accuracy: The ability of a recommender system to predict a user's rating for an item
    - Mean absolute error (MAE)
  - Rank accuracy
    - Precision: percentage of items in a recommendation list that the user would rate as useful
    - Half-life utility: the utility of a recommendation list to a user
      - Assumption: the likelihood that a user examines a recommended object decays exponentially with the objects ranking

# Practical Issues: Ratings

- Explicit ratings:
  - Users rate themselves for an item
  - Most accurate descriptions of a user's preference
  - Challenging in collecting data
- Implicit ratings
  - Observations of user behavior
  - Can be collected with little or no cost to user
  - Ratings inference may be imprecise.

The image displays two side-by-side screenshots of an Amazon.com product page for the book "A Little Book on Perl".

**Screenshot 1 (Left): Rating and Similar Items**

This screenshot shows the main product page. At the bottom, there is a red-bordered box containing the heading "This Book and You" and a "Rate this item" section. The "Rate this item" section includes a 5-star rating input field with the number "1" highlighted, and a checkbox labeled "Not interested". Below this are links: "Write a Review", "Write a So You'd Like To... Guide", and "E-mail a Friend About This Item".

**Screenshot 2 (Right): Recommendations**

This screenshot shows the "So You'd Like To..." section, which lists recommended items based on the user's purchase history. It includes titles like "Learn to Program", "design a database driven web application or website!", "Make a Website!", and "Create a So You'd Like to... guide". Below this is the "Look for similar items by category" section, which lists categories under "Books > Computers & Internet > Books > Programming > Perl".

## Practical Issues: Cold start

Not enough data in order to make accurate recommendations

- New user
  - Rate some initial items
  - Non-personalized recommendations - Describe tastes
  - Demographic info.
- New Item
  - Non-CF : content analysis, metadata

# Practical Issues: Others

- Scalability:
  - Usually, there are millions of users and products
  - Large amount of computation power is necessary
- Sparsity:
  - The number of items is extremely large
  - The most active users will only have rated a small subset of the overall database
  - Even the most popular items have a very few ratings
- Privacy and Trust:
  - Users profiles, Personalized information
  - Recommender system may break trust when malicious users give ratings that are not representative of their true preferences.

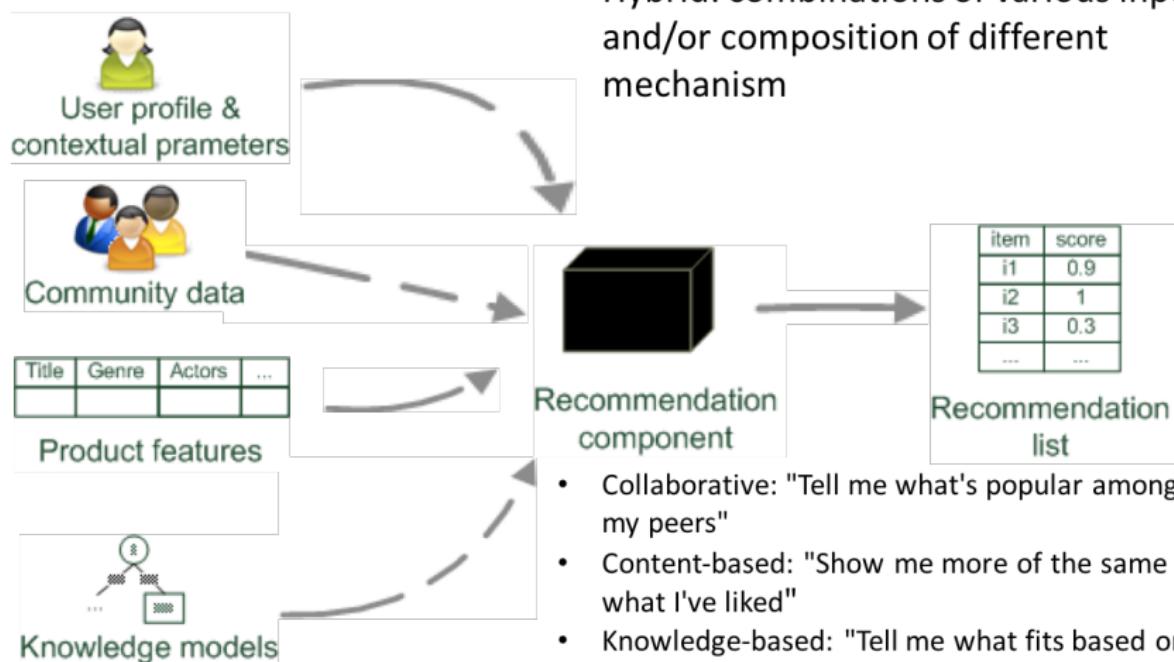
# Collaborative Filtering: Summary

- Pros: well-understood, works well in some domains, no knowledge engineering required
- Cons: requires user community, sparsity problems, no integration of other knowledge sources, no explanation of results
- What is the best CF method? Differences between methods could be very small with limited data
- How to evaluate the prediction quality? MAE / RMSE
  - What does an MAE of 0.7 actually mean?

# Outline

- 1 Administration
- 2 Recommender Systems: Introduction
- 3 Content-based recommendations
- 4 Collaborative Filtering
- 5 Hybrids

# Hybrid recommendation



# Credits

Parts of this lecture note is based on the slides of:

- Prof. Yan Liu
- Narges Razavian@ csail.mit
- Alex Smola@CMU
- Yahuda Koren@Yahoo labs
- Bing Liu@UIC
- Xaviar Amatriain@ Netflix
- Parisa Mansourifard